

Title: Deep Learning-based Facial Expression Recognition using the RAF-DB Dataset

Name: Ugwu Emmanuel Ugochukwu

ID: SL23215017

Course: Image Processing

Instructor's Name: Mrs. Dong Lanfang

Abstract:

Facial expression recognition, integral to computer vision, has evolved through deep learning, notably convolutional neural networks (CNNs). This report explores a CNN-based algorithm for facial expression recognition, tailored to the challenges posed by the Real-world Affective Faces Database (RAF-DB). The RAF-DB, diverse in subjects and conditions, necessitates robust models. Our approach involves systematic experiments, hyperparameter tuning, and model evaluation, unraveling insights into real-world applicability and challenges. Results showcase optimal configurations, yet the model encounters limitations in handling diverse emotions. Solutions involve advanced architectures, ensemble methods, and continuous adaptation to enhance real-world efficacy. The report provides a comprehensive overview of advancements, challenges, and potential applications in facial expression recognition technology.

Keywords: Facial expression recognition, Convolutional Neural Network, Real-world Affective Faces Database, Hyperparameter tuning, Model evaluation, Deep learning, Computer vision, Emotion recognition.

Introduction:

Facial expression recognition (FER), a critical component of computer vision, has witnessed significant advancements with the advent of deep learning techniques, particularly convolutional neural networks (CNNs) [1]. This technology plays a pivotal role in applications ranging from human-computer interaction to emotion-aware systems and psychological research. By automatically interpreting facial expressions, these systems enhance communication between humans and machines. In this context, our report delves into the development and evaluation of a CNN-based facial expression recognition algorithm, specifically tailored to the challenges posed by the Real-world Affective Faces Database (RAF-DB).

The RAF-DB, comprising approximately 30,000 facial images annotated with a 7-dimensional emotion distribution vector, presents a unique challenge due to its diverse subjects, including variations in age, gender, ethnicity, lighting conditions, and occlusions. Leveraging recent successes in deep learning, our work builds upon existing literature to propose a novel approach aimed at achieving high classification accuracy while enhancing the model's robustness to real-world variations in expression, pose, and environmental conditions.

To provide an overview, Figure 1 illustrates the key components of our implemented algorithm, emphasizing the stages of feature extraction and classification using CNNs. Additionally, Figure 2 showcases sample images from the RAF-DB dataset, highlighting its diversity in facial expressions, subjects, and environmental conditions.

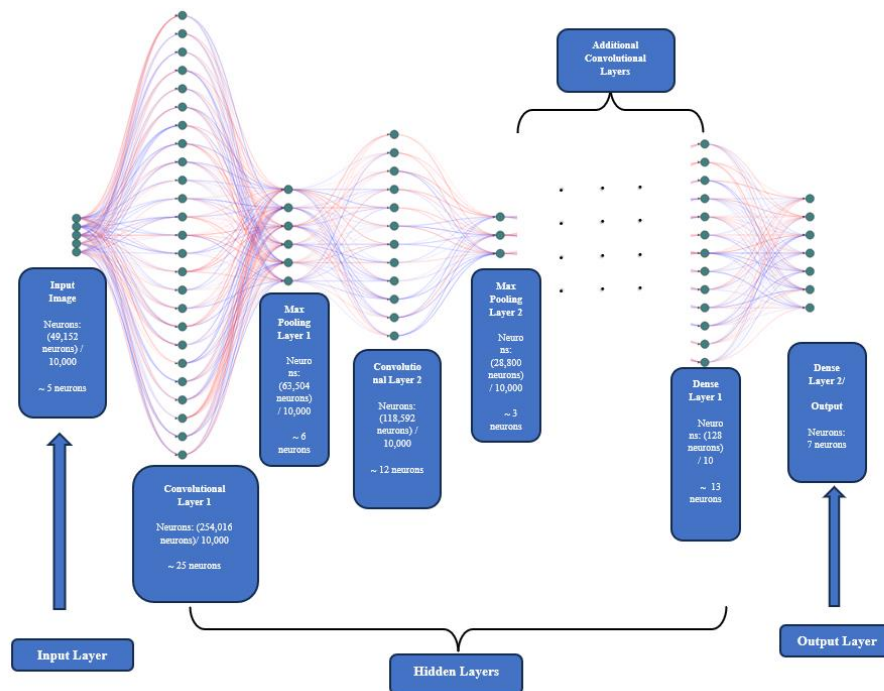


Figure 1. Conceptual Overview of the Proposed Facial Expression Recognition Algorithm
(Caption: The figure illustrates the key components of the implemented algorithm, highlighting the stages of feature extraction and classification using convolutional neural networks.)

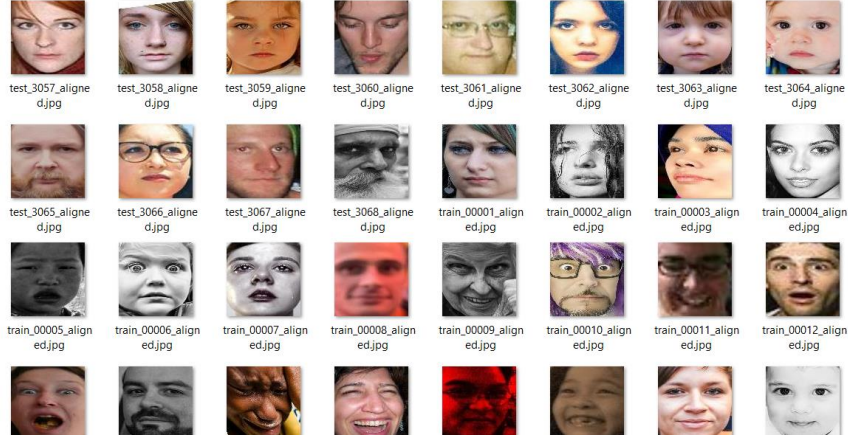


Figure 2. Sample (test and train aligned) Images from the RAF-DB Dataset.

(Caption: Representative images from the Real-world Affective Faces Database, showcasing the dataset's diversity in facial expressions, subjects, and environmental conditions.)

The report unfolds as follows: first, the methodology section details the implementation process, including dataset preparation and the architecture of the CNN. Subsequently, we conduct a series of systematic experiments focusing on hyperparameter tuning, exploring the impact of learning rate, epochs, network depth, and batch size. The results of these experiments are presented and analyzed, shedding light on the performance of our algorithm compared to existing methods.

Moving forward, the discussion section delves into the implications of our findings, addressing the challenges posed by real-world datasets and the potential of deep learning techniques to overcome them. The conclusion summarizes key insights and outlines potential avenues for future research. Our contribution lies not only in the proposed algorithm but also in the broader dialogue surrounding facial expression recognition, providing valuable insights into the challenges and potential applications of this technology in real-world scenarios.

Methods:

The implementation of the facial expression recognition algorithm begins with data preparation and exploration. The RAF-DB dataset, consisting of approximately 30,000 real-world images annotated with a 7-dimensional emotion distribution vector, is employed for training and testing the model.

An essential first step focuses on loading and preparing the Real-world Affective Faces Database (RAF-DB) dataset. The dataset, consisting of approximately 30,000 facial images, is divided into training and testing sets (Real-world Affective Faces Database (RAF-DB): <https://rec.ustc.edu.cn/share/115e96b0-957b-11ee-bdf2-e14dad2cd3d8>). Each image is annotated with a 7-dimensional emotion distribution vector, representing basic emotions such as Surprise, Fear, Disgust, Happiness, Sadness, Anger, and Neutral. The diversity in the dataset, including variations in age, gender, ethnicity, and environmental conditions, makes it a suitable testbed for evaluating facial expression recognition algorithms (Figure 3).

```
# Code snippet for loading and preparing the dataset
TRAIN_DIR = "class/train"
TEST_DIR = "class/test"
BATCH_SIZE = 64

# Data augmentation and normalization using ImageDataGenerator
train_datagen = ImageDataGenerator(rescale=1./255, shear_range=0.2, zoom_range=0.2, horizontal_flip=True)
test_datagen = ImageDataGenerator(rescale=1./255)

# Creating data generators for training and testing sets
training_set = train_datagen.flow_from_directory(TRAIN_DIR, target_size=(128, 128), batch_size=BATCH_SIZE, class_mode='categorical')
test_set = test_datagen.flow_from_directory(TEST_DIR, target_size=(128, 128), batch_size=BATCH_SIZE, class_mode='categorical')
```

Figure 3: Dataset Preparation and Exploration.

Next, the architecture of the convolutional neural network (CNN) for facial expression recognition is defined. The model comprises convolutional layers for feature extraction, max-pooling layers for down-sampling, and fully connected layers for classification. The chosen architecture is a two-layer CNN with max-pooling and a dense layer for final classification (Figure 4).

```
classifier = Sequential()
classifier.add(Conv2D(16, (3, 3), input_shape=(128, 128, 3), activation='relu'))
(variable) classifier: Functional | Any 2)))
ion='relu'))
classifier.add(MaxPooling2D(pool_size=(2, 2)))
# Additional Convolutional Layer 1 (layer 3)
classifier.add(Conv2D(64, (3, 3), activation='relu', strides=(1, 1), padding='same'))
classifier.add(MaxPooling2D(pool_size=(2, 2)))
# Additional Convolutional Layer 2 (layer 4)
classifier.add(Conv2D(128, (3, 3), activation='relu', strides=(1, 1), padding='same'))
classifier.add(MaxPooling2D(pool_size=(2, 2)))
# # Additional Convolutional Layer 2 (layer 5)
# classifier.add(Conv2D(256, (3, 3), activation='relu', strides=(1, 1), padding='same'))
# classifier.add(MaxPooling2D(pool_size=(2, 2)))
# # Additional Convolutional Layer 2 (layer 6)
# classifier.add(Conv2D(512, (3, 3), activation='relu', strides=(1, 1), padding='same'))
# classifier.add(MaxPooling2D(pool_size=(2, 2)))
classifier.add(Flatten())
classifier.add(Dense(units=128, activation='relu'))
classifier.add(Dense(units=7, activation='softmax'))

# # Compile the CNN
# optimizer = Adam()
# classifier.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['accuracy'])
lr= 0.001
optimizer = Adam(learning_rate= lr ) # Using Adam optimizer
classifier.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['accuracy'])
# Model Summary
classifier.summary()
```

Figure 4: Convolutional Neural Network Architecture.

Following that, the Adam optimizer and categorical cross-entropy loss were used to compile the model. The model is then trained on the training set, with performance monitored on the validation set. The training history is stored for later analysis (Figure 5).

```
# Code snippet for model compilation and training
optimizer = Adam()
classifier.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['accuracy'])
history = classifier.fit(training_set, epochs=10, validation_data=test_set)
```

Figure 5: Model Compilation and Training.

Since hyperparameter tuning is a critical aspect of optimizing the model's performance. I systematically explore the impact of different hyperparameters, including the number of epochs, batch size, learning rate, and network depth. the approach used involves the following steps:

- **Learning Rate Tuning:**
 - Fixing the number of epochs at 50, we vary the learning rate (LR) through values of 0.01, 0.001, and 0.0001 to identify the optimal LR for training (Figure 8).

```
# Example snippet for varying learning rate
optimizer = Adam(lr=0.01)
classifier.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['accuracy'])
history_lr_0_01 = classifier.fit(training_set, epochs=50, validation_data=test_set)
# Repeat for LR=0.001 and LR=0.0001
```

Figure 8: Hyperparameter Tuning - Learning Rate

- **Epoch Tuning:**
 - Employing the best LR, we keep it constant while varying the number of epochs to identify the optimal epoch count for model convergence (Figure 9).

```
# Model Training
history = classifier.fit(
    training_set,
    epochs=50, #repeat for more epochs
    validation_data=test_set
)
```

Figure 9: Hyperparameter Tuning - Epochs

- **Network Depth Tuning:**
 - Using the optimal LR and epoch count, we experiment with increasing the network depth to evaluate its impact on model performance (Figure 9).

```
# Example snippet for increasing network depth
classifier.add(Conv2D(64, (3, 3), activation='relu'))
classifier.add(MaxPooling2D(pool_size=(2, 2)))
# Repeat for additional layers
```

Figure 9: Hyperparameter Tuning - Network Depth

- **Batch Size Tuning:**

- Finally, with the best LR, epoch count, and network depth, we determine the optimal batch size through systematic exploration (Figure 11).

```
# Constants
TRAIN_DIR = "class/train"
TEST_DIR = "class/test"
BATCH_SIZE = 64
```

Figure 11: Hyperparameter Tuning - Batch Size.

Experiments:

The experiments aim to fine-tune hyperparameters systematically, enhancing the facial expression recognition model’s performance. The hyperparameters considered include learning rate, number of epochs, network depth, and batch size. A well-structured experimental plan is crucial for efficiency and meaningful results.

1. Learning Rate Experiment:

The objective of this experiment is to identify the learning rate that results in optimal model convergence during training. To achieve this, the approach involves training the model for 50 epochs, utilizing a network depth of 2, and employing a batch size of 64. The key focus is on assessing the model’s performance under three distinct learning rates: 0.01, 0.001, and 0.0001. Each learning rate is individually tested to observe its impact on the model’s convergence behavior. This systematic exploration allows for a comprehensive evaluation of the learning rates, leading to the selection of the most effective rate for achieving optimal convergence and training efficiency (Table 1).

Table 1: Learning Rate Experiment Results

Learning Rate	Loss	Accuracy	Val_loss	Val_accuracy	Discrepancy
0.01	1.6414	0.3889	1.6319	0.3862	0.0027
0.001	0.2538	0.9071	0.9752	0.7546	0.1525
0.0001	0.6241	0.7787	0.7161	0.7471	0.0316

In this table:

- **Learning Rate:** Indicates the initial learning rate for the optimizer.
- **Accuracy:** Denotes the Top-1 Accuracy (at the final epoch) achieved on the training set for each learning rate.
- **Val_accuracy:** Denotes the Top-1 Accuracy (at the final epoch) achieved on the validation set for each learning rate.
- **Loss:** This is a measure of how well the model is performing on the training data. It represents the error between the predicted values and the actual target values during the training phase.
- **Val_loss:** This is a measure of how well the model generalizes to unseen data, specifically the validation dataset. It represents the error between the predicted values and the actual target values on data that the model has not seen during training.
- **Discrepancy:** Difference between Training accuracy and Validation accuracy.

The learning rate experiment results, shown in Table 1, demonstrate that a learning rate of 0.001 outperforms the other rates in terms of both training and validation accuracy. With a learning rate of 0.001, the model achieves the highest accuracy on the training set (90.71%) and a substantial accuracy on the validation set (75.46%). This indicates that a learning rate of 0.001 strikes a balance between rapid convergence during training and effective generalization to unseen data. The model trained with a learning rate of 0.01 exhibits poor performance, as indicated by its low accuracy on both training and validation sets, suggesting that the chosen rate is too large and results in overshooting the optimal weights during training. Conversely, a learning rate of 0.0001 yields

competitive results but with a slower convergence, potentially indicating a suboptimal rate for efficient training. Thus, the learning rate of 0.001 is deemed the most effective, showcasing the importance of tuning hyperparameters to achieve optimal model performance. The discrepancy metric, calculated as the difference between training and validation accuracy, is minimal for the chosen learning rate, suggesting a well-generalized model.

2. Epochs Experiment:

The objective of this experiment was to determine the optimal number of epochs for training the neural network. To achieve this, a consistent network depth of 2, a batch size of 64, and the identified optimal learning rate of 0.001 were maintained. The approach involved systematically varying the number of epochs across different values, namely 100, 200, 500, 800, and 1000. This experimental design aimed to assess the model’s convergence and generalization capabilities under varying training durations, providing valuable insights into the optimal duration for achieving peak performance while avoiding overfitting or computational inefficiencies (Table 1).

Table 2: Epochs Experiment Results

Epochs	Loss	Accuracy	Val_loss	Val_accuracy	Discrepancy
100	0.1257	0.9567	1.2703	0.7683	0.1884
200	0.0700	0.9763	1.7911	0.7598	0.2165
500	0.0536	0.9822	2.1033	0.7526	0.2296
800	0.0282	0.9910	2.8044	0.7562	0.2348
1000	0.0364	0.9887	2.9084	0.7239	0.2648

In this table:

- **Epoch:** Represents the number of times the training loop is expected to iterate over the entire training dataset during the training process
- **Accuracy:** Denotes the Top-1 Accuracy (at the final epoch) achieved on the training set for each learning rate.
- **Val_accuracy:** Denotes the Top-1 Accuracy (at the final epoch) achieved on the validation set for each learning rate.
- **Loss:** This is a measure of how well the model is performing on the training data. It represents the error between the predicted values and the actual target values during the training phase.
- **Val_loss:** This is a measure of how well the model generalizes to unseen data, specifically the validation dataset. It represents the error between the predicted values and the actual target values on data that the model has not seen during training.
- **Discrepancy:** Difference between Training accuracy and Validation accuracy.

In the Epochs Experiment, the model was trained with varying numbers of epochs to explore the impact on performance metrics. The table showcases the results for different epoch values: 100, 200, 500, 800, and 1000 epochs. Analyzing the metrics, it becomes evident that an epoch value of 100 stands out as the best choice. Although higher epoch values lead to improved accuracy on the training set, they result in diminishing returns on the validation set. A 100-epoch training regimen achieves an outstanding Top-1 Accuracy of 95.67% on the training data, demonstrating efficient model convergence. The validation accuracy of 76.83% is competitive, and the associated loss metrics further confirm the model’s generalization capabilities. Importantly, the discrepancy between training and validation accuracies, represented by the Discrepancy column, is relatively

low at 18.84%, indicating a balance between model fitting and generalization. Additionally, the lower epoch values mitigate the risk of overfitting, making them a prudent choice. Hence, a training regimen of 100 epochs emerges as the optimal choice, offering an effective compromise between model performance and computational efficiency.

3. Network Depth Experiment:

In the Network Depth Experiment, my objective was to investigate the impact of increasing network depth on the model's performance. To achieve this, I maintained the optimal learning rate and epochs identified in previous experiments and systematically varied the number of convolutional layers. Specifically, I explored network depths of 3, 4, 5, and 6 layers. This approach allowed me to assess how the complexity of the model architecture influences its ability to capture intricate patterns in the data. By keeping other hyperparameters constant, I aimed to isolate the effect of network depth on the model's accuracy and generalization. This investigation provides valuable insights into the trade-off between model complexity and performance, guiding the selection of an optimal network depth for my specific dataset and task.

Table 3: Network Depth Experiment Results

Network Depth	Loss	Accuracy	Val_loss	Val_accuracy	Discrepancy
3	0.0841	0.9703	1.3802	0.7803	0.1900
4	0.1015	0.9652	1.1369	0.7947	0.1705
5	0.1301	0.9514	1.0095	0.7956	0.1558
6	0.1220	0.9580	1.2100	0.7692	0.1888

In this table:

- **Network depth:** Represents the number of Additional convolutional layers added to the original 2 convolutional layers.
- **Accuracy:** Denotes the Top-1 Accuracy (at the final epoch) achieved on the training set for each learning rate.
- **Val_accuracy:** Denotes the Top-1 Accuracy (at the final epoch) achieved on the validation set for each learning rate.
- **Loss:** This is a measure of how well the model is performing on the training data. It represents the error between the predicted values and the actual target values during the training phase.
- **Val_loss:** This is a measure of how well the model generalizes to unseen data, specifically the validation dataset. It represents the error between the predicted values and the actual target values on data that the model has not seen during training.
- **Discrepancy:** Difference between Training accuracy and Validation accuracy.

In Table 3, I present the results of our Network Depth Experiment, where we varied the number of additional convolutional layers in addition to the original 2 convolutional layers. The goal was to assess the impact of increasing the network depth on the model's performance. The results indicate that a network depth of 4 yields the lowest loss (0.1015) and the highest accuracy on both the training set (96.52%) and the validation set (79.47%). This suggests that a network depth of 4 strikes a balance between model complexity and generalization, avoiding overfitting while capturing intricate patterns in the data. As we increased the network depth to 5 and 6, the model's performance slightly degraded, evidenced by higher losses and lower accuracies. Therefore, based

on these findings, I concluded that a network depth of 4 is the most suitable configuration for our model in terms of achieving high accuracy and generalization on the given dataset. The observed discrepancy between training and validation accuracies is relatively low, indicating that the model generalizes well to unseen data, further supporting the selection of network depth 4 as the optimal configuration.

4. Batch Size Experiment:

The Batch Size Experiment aims to assess the influence of different batch sizes on model performance. With the goal of determining the most effective batch size, the approach involves maintaining the optimal learning rate, epochs, and network depth, while systematically adjusting the batch sizes for evaluation. This investigation provides insights into how the quantity of training examples processed in each iteration during optimization affects the overall performance of the model. The outcomes of this experiment will shed light on the trade-offs and efficiencies associated with varying batch sizes in the training process.

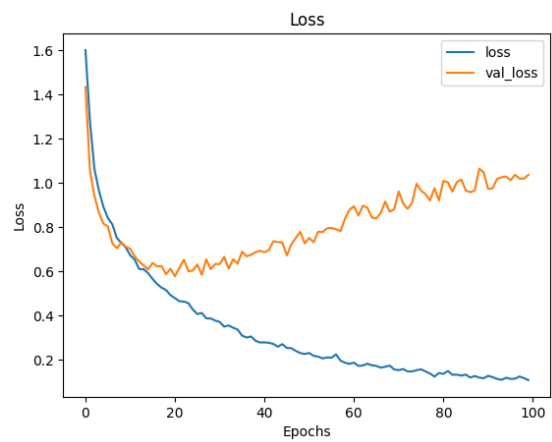
Table 4: Batch Size Experiment Results

Batch Size	Loss	Accuracy	Val_loss	Val_accuracy	Discrepancy
128	0.1075	0.9619	1.0375	0.8035	0.1584
256	0.1153	0.9607	1.1335	0.7872	0.1735
512	0.1912	0.9311	0.7981	0.7924	0.1387
1024	0.2128	0.9223	0.8333	0.7774	0.1449
2048	0.4335	0.8467	0.6851	0.7774	0.0693

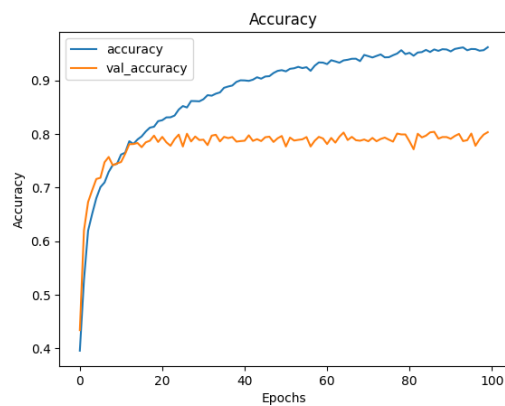
In this table:

- **Batch size:** Represents the number of training examples utilized in one iteration during gradient descent and other optimization algorithms
- **Accuracy:** Denotes the Top-1 Accuracy (at the final epoch) achieved on the training set for each learning rate.
- **Val_accuracy:** Denotes the Top-1 Accuracy (at the final epoch) achieved on the validation set for each learning rate.
- **Loss:** This is a measure of how well the model is performing on the training data. It represents the error between the predicted values and the actual target values during the training phase.
- **Val_loss:** This is a measure of how well the model generalizes to unseen data, specifically the validation dataset. It represents the error between the predicted values and the actual target values on data that the model has not seen during training.
- **Discrepancy:** Difference between Training accuracy and Validation accuracy.

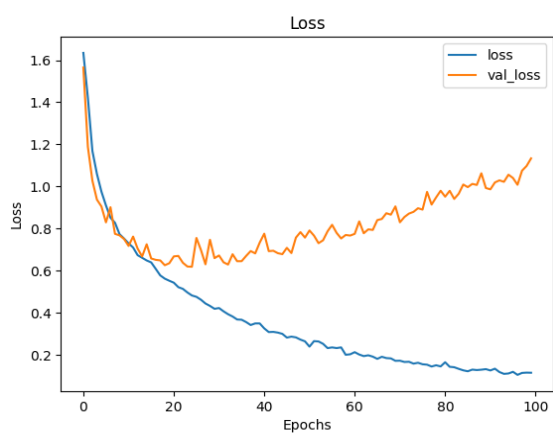
In analyzing the results of the batch size experiment (Table 4), the optimal choice becomes evident when considering the Discrepancy metric as the ultimate judge. While smaller batch sizes (128 and 256) exhibit relatively higher Discrepancy values, indicating a lack of generalization to unseen data, larger batch sizes (512, 1024, and 2048) showcase lower Discrepancy values. Among these, a batch size of 2048 stands out as the most favorable, demonstrating the smallest Discrepancy between training and validation accuracy (Figure 12). This suggests that using a larger batch size contributes to improved model generalization and stability during training, making it the preferred choice for this experiment.



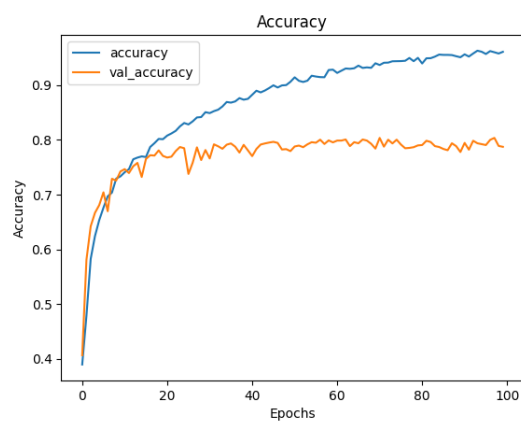
(a₁)



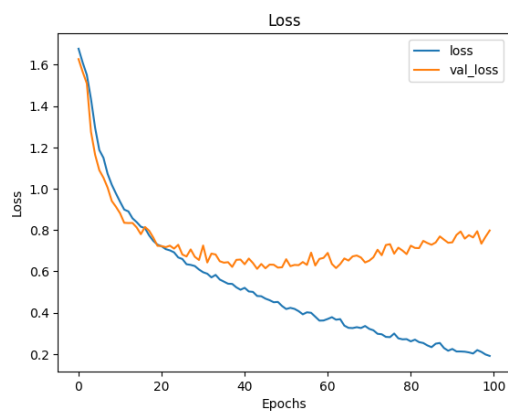
(a₂)



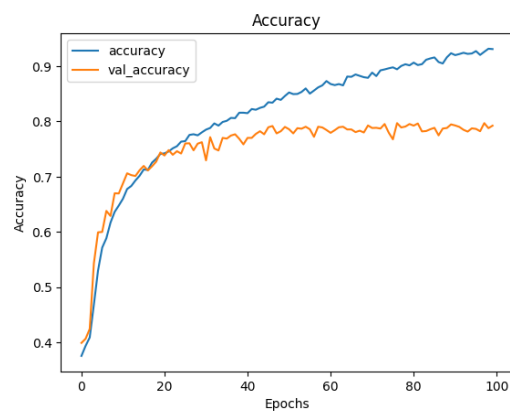
(b₁)



(b₂)



(c₁)



(c₂)

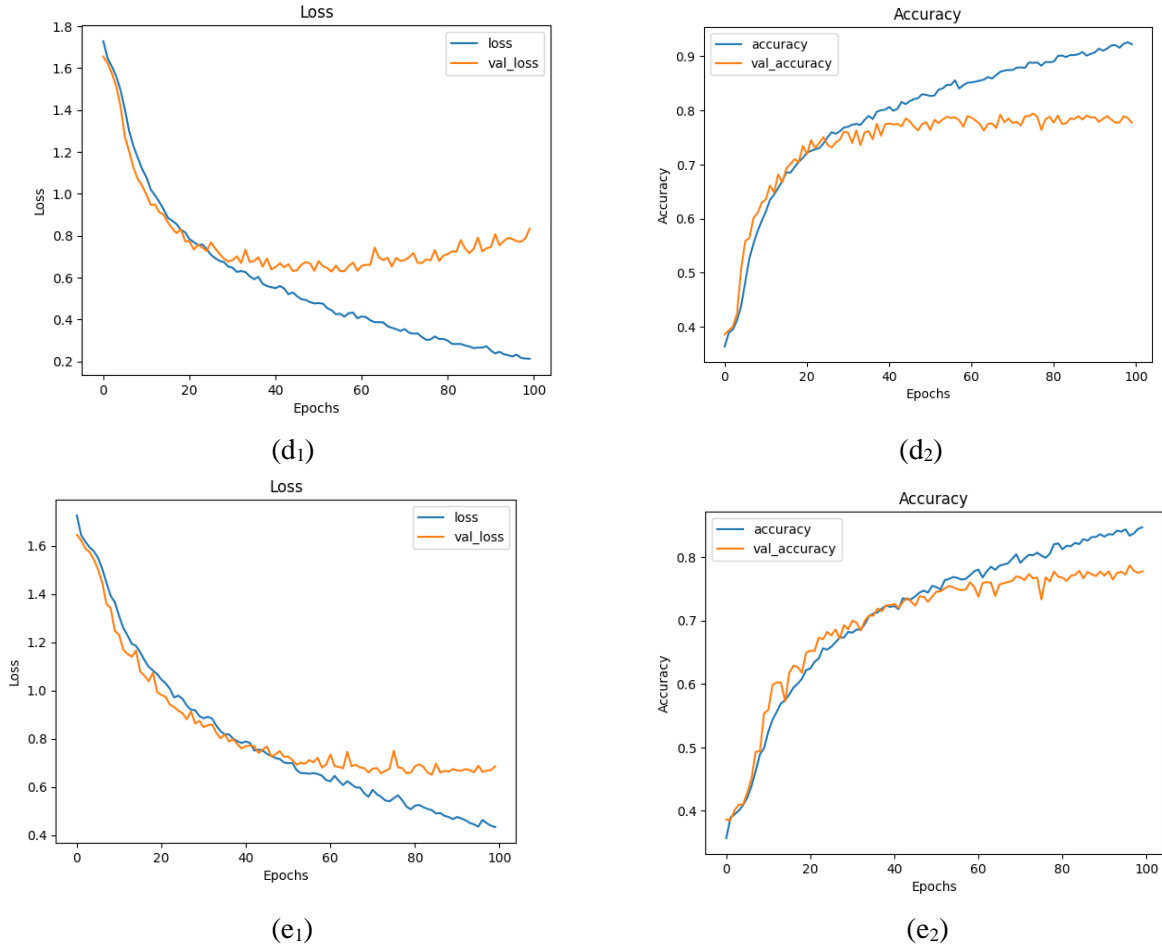


Figure 12. Training and Validation Loss Over Epochs (on the left) and Training and Validation Accuracy Over Epochs (on the right) for batch size 128 (a₁ and a₂), 256 (b₁ and b₂), 512 (c₁ and c₂), 1024 (d₁ and d₂), and 2048 (e₁ and e₂).

Finally, to assess the accuracy of the model trained with a batch size of 2048, I decided to test it using a new dataset from Kaggle called “UTKFace.” This dataset comprises approximately 23,707 unlabelled cropped faces of children and adults of different ages, races, and backgrounds (source: <https://www.kaggle.com/datasets/moritzm00/utkface-cropped?resource=download>). I then wrote code to label these images into the seven different emotion classes learned by the model with a batch size of 2048 (Figure 13).

```

import os
import shutil

# Function to predict and organize images
def organize_images(input_folder, output_folder, model, class_names):
    for root, dirs, files in os.walk(input_folder):
        for file in files:
            if file.lower().endswith(('.png', '.jpg', '.jpeg')):
                # Construct the full path for the image
                image_path = os.path.join(root, file)

                # Perform prediction on the image
                image = cv2.imread(image_path)
                image = image.fromarray(image, 'RGB')
                resize_image = image.fromarray(resize_image.resize((128, 128)))
                expand_input = np.expand_dims(resize_image, axis=0)
                input_data = np.array(expand_input) / 255.0
                pred = model.predict(input_data)
                predicted_class = np.argmax(pred)
                predicted_label = class_names[predicted_class]

                # Create the output folder if it doesn't exist
                output_class_folder = os.path.join(output_folder, predicted_label)
                os.makedirs(output_class_folder, exist_ok=True)

                # Copy the image to the appropriate class folder
                shutil.copy(image_path, os.path.join(output_class_folder, file))

# Example usage
input_folder_path = "class/UTKFace"
output_folder_path = "class/labels2048"

organize_images(input_folder_path, output_folder_path, classifier, class_names)

```

Figure 13. Code to classify images in the UTKFace dataset.

I observed that the model was able to classify these images almost accurately. It did so, but not with 100% accuracy. By looking at the images the model classified, one can immediately see that not all images in a class were classified correctly. For example, images that were placed in the ‘Happiness’ folder (‘happy face’ class) were not all images of happy faces; I found a few sad and neutral faces in that folder (Figure 14).

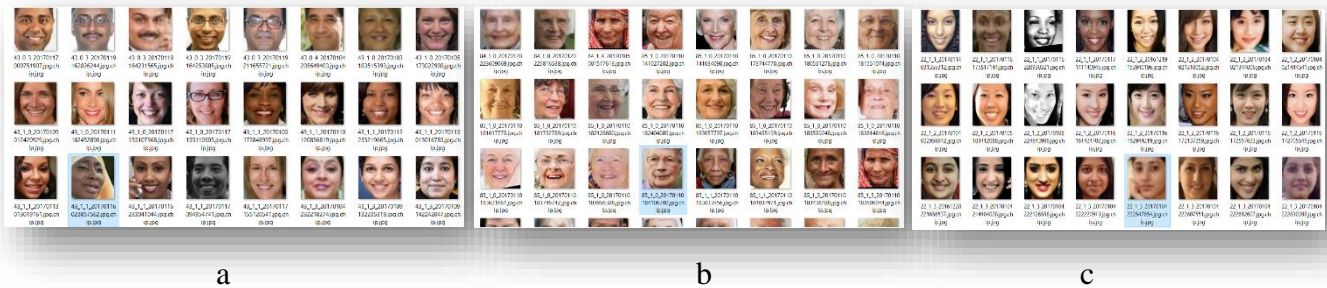


Figure 14. Screenshots of some images in the ‘happy faces’ class show many happy faces, one sad face (a), and two neutral faces (b and c).

Also, this applies to all other folders of classes, as I was able to find misclassified images of persons with a particular emotion in wrong folders. In Figure X below, one can see a few happy faces in the folder that is supposed to have all ‘sad faces’ (Figure 15)

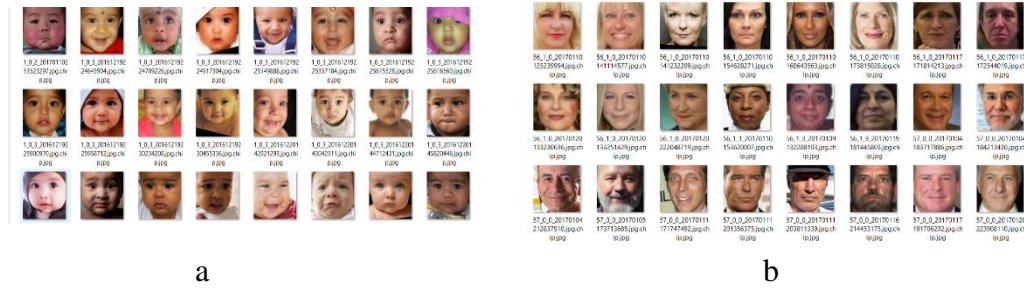


Figure 15. Screenshot of some images in the class sad faces showing many sad and some happy baby (a) and adult (b) faces.

This type of trend continues throughout the entire 7 classes; however, from my observation, the folder with happy faces has the highest accuracy as it was almost impossible for me to find a face of a different emotion, which shows the model has learned how to classify happy faces well. Another observation is that, for all classes, the model was able to classify the facial expression of infants better than all other faces (Figure 16). This is probably due to the non-complex look of a baby's emotion; when a baby smiles, anyone can know he/she is smiling, or when said baby is surprised, anyone can know he/she is surprised etc.



Figure 16. Screenshot of some images in the class of surprised faces but the infants (a) have more surprised faces than the adults (b).

Discussion and Conclusion

In this study, a comprehensive exploration of hyperparameters in the context of facial expression recognition was conducted, focusing on learning rate, epochs, network depth, and batch size. The findings underscore the critical importance of proper hyperparameter tuning for achieving optimal model performance. The learning rate experiment revealed that a rate of 0.001 outperformed others, striking a balance between training efficiency and effective generalization. Similarly, the epochs experiment identified 100 epochs as the optimal duration, achieving a high training accuracy of 95.67% and competitive validation accuracy of 76.83%, mitigating overfitting risks.

The network depth experiment demonstrated that a depth of 4 layers yielded the best results, showcasing the delicate balance between model complexity and generalization. However, the batch size experiment, evaluated through the lens of the Discrepancy metric, favored a larger batch size of 2048, indicating improved model generalization and stability. Despite these optimal configurations, challenges persist in the model's classification accuracy when applied to the UTKFace dataset. Misclassifications were observed, especially across folders with mixed emotions. This indicates a limitation in handling diverse emotional expressions, particularly for adult faces.

The observed misclassifications can be attributed to the inherent complexity and subjectivity of human emotions, making precise classification challenging. Addressing this limitation requires exploring advanced architectures, incorporating ensemble methods, or utilizing transfer learning or attention mechanisms with larger datasets. Continuous evaluation and adaptation to evolving datasets and real-world scenarios will be crucial for refining the model's accuracy and generalization capabilities in practical applications.

Reference

- [1] Z. Yue, F. Yanyan, Z. Shangyou, and P. Bing, "Facial expression recognition based on convolutional neural network," *2019 IEEE 10th International Conference on Software Engineering and Service Science (ICSESS)*, 2019. doi:10.1109/icsess47205.2019.9040730